

Home Credit Default Risk

PySpark Introduction

- Apache Spark is popular in big data due to in-memory computation and parallel processing.
- MLlib, built on Spark, is a scalable Machine Learning library with high-quality algorithms and speed.
- MLlib has APIs for Java, Python, and Scala, making it suitable for Data Analysts, Engineers, and Scientists.
- MLlib includes algorithms for classification, regression, clustering, collaborative filtering, and more.

Problem Statement Introduction

- The article discusses building an end-to-end machine learning model using MLlib in PySpark.
- The dataset used is from the Home Credit Default Risk competition on Kaggle.
- The objective is to determine if loan applicants can repay their loans based on collected data.
- It's a binary classification problem with an imbalanced target label: 0 (applicants who paid back loans) and 1 (applicants who didn't).
- The distribution ratio is approximately 0.91 (applicants who repaid) to 0.09 (applicants who didn't).

1. Data Ingestion and Spark session creation

- Dataset Link: <https://www.kaggle.com/c/home-credit-default-risk> (<https://www.kaggle.com/c/home-credit-default-risk>)

```
In [ ]: from pyspark.sql import SparkSession
        # initiate our session and read the main CSV file, then we print the #dataframe schema
        spark = SparkSession.builder.appName('imbalanced_binary_classification').getOrCreate()
        new_df = spark.read.csv("../Data/application_train.csv/application_train.csv", header=True, inferSchema=True)
        new_df.printSchema()
```

root

```
|-- SK_ID_CURR: integer (nullable = true)
|-- TARGET: integer (nullable = true)
|-- NAME_CONTRACT_TYPE: string (nullable = true)
|-- CODE_GENDER: string (nullable = true)
|-- FLAG_OWN_CAR: string (nullable = true)
|-- FLAG_OWN_REALTY: string (nullable = true)
|-- CNT_CHILDREN: integer (nullable = true)
|-- AMT_INCOME_TOTAL: double (nullable = true)
|-- AMT_CREDIT: double (nullable = true)
|-- AMT_ANNUITY: double (nullable = true)
|-- AMT_GOODS_PRICE: double (nullable = true)
|-- NAME_TYPE_SUITE: string (nullable = true)
|-- NAME_INCOME_TYPE: string (nullable = true)
|-- NAME_EDUCATION_TYPE: string (nullable = true)
|-- NAME_FAMILY_STATUS: string (nullable = true)
|-- NAME_HOUSING_TYPE: string (nullable = true)
|-- REGION_POPULATION_RELATIVE: double (nullable = true)
|-- DAYS_BIRTH: integer (nullable = true)
|-- DAYS_EMPLOYED: integer (nullable = true)
|-- DAYS_REGISTRATION: double (nullable = true)
|-- DAYS_ID_PUBLISH: integer (nullable = true)
|-- OWN_CAR_AGE: double (nullable = true)
|-- FLAG_MOBIL: integer (nullable = true)
|-- FLAG_EMP_PHONE: integer (nullable = true)
|-- FLAG_WORK_PHONE: integer (nullable = true)
|-- FLAG_CONT_MOBILE: integer (nullable = true)
|-- FLAG_PHONE: integer (nullable = true)
|-- FLAG_EMAIL: integer (nullable = true)
|-- OCCUPATION_TYPE: string (nullable = true)
|-- CNT_FAM_MEMBERS: double (nullable = true)
|-- REGION_RATING_CLIENT: integer (nullable = true)
|-- REGION_RATING_CLIENT_W_CITY: integer (nullable = true)
|-- WEEKDAY_APPR_PROCESS_START: string (nullable = true)
|-- HOUR_APPR_PROCESS_START: integer (nullable = true)
|-- REG_REGION_NOT_LIVE_REGION: integer (nullable = true)
|-- REG_REGION_NOT_WORK_REGION: integer (nullable = true)
|-- LIVE_REGION_NOT_WORK_REGION: integer (nullable = true)
|-- REG_CITY_NOT_LIVE_CITY: integer (nullable = true)
|-- REG_CITY_NOT_WORK_CITY: integer (nullable = true)
|-- LIVE_CITY_NOT_WORK_CITY: integer (nullable = true)
```

```
|-- ORGANIZATION_TYPE: string (nullable = true)
|-- EXT_SOURCE_1: double (nullable = true)
|-- EXT_SOURCE_2: double (nullable = true)
|-- EXT_SOURCE_3: double (nullable = true)
|-- APARTMENTS_AVG: double (nullable = true)
|-- BASEMENTAREA_AVG: double (nullable = true)
|-- YEARS_BEGINEXPLUATATION_AVG: double (nullable = true)
|-- YEARS_BUILD_AVG: double (nullable = true)
|-- COMMONAREA_AVG: double (nullable = true)
|-- ELEVATORS_AVG: double (nullable = true)
|-- ENTRANCES_AVG: double (nullable = true)
|-- FLOORSMAX_AVG: double (nullable = true)
|-- FLOORSMIN_AVG: double (nullable = true)
|-- LANDAREA_AVG: double (nullable = true)
|-- LIVINGAPARTMENTS_AVG: double (nullable = true)
|-- LIVINGAREA_AVG: double (nullable = true)
|-- NONLIVINGAPARTMENTS_AVG: double (nullable = true)
|-- NONLIVINGAREA_AVG: double (nullable = true)
|-- APARTMENTS_MODE: double (nullable = true)
|-- BASEMENTAREA_MODE: double (nullable = true)
|-- YEARS_BEGINEXPLUATATION_MODE: double (nullable = true)
|-- YEARS_BUILD_MODE: double (nullable = true)
|-- COMMONAREA_MODE: double (nullable = true)
|-- ELEVATORS_MODE: double (nullable = true)
|-- ENTRANCES_MODE: double (nullable = true)
|-- FLOORSMAX_MODE: double (nullable = true)
|-- FLOORSMIN_MODE: double (nullable = true)
|-- LANDAREA_MODE: double (nullable = true)
|-- LIVINGAPARTMENTS_MODE: double (nullable = true)
|-- LIVINGAREA_MODE: double (nullable = true)
|-- NONLIVINGAPARTMENTS_MODE: double (nullable = true)
|-- NONLIVINGAREA_MODE: double (nullable = true)
|-- APARTMENTS_MEDI: double (nullable = true)
|-- BASEMENTAREA_MEDI: double (nullable = true)
|-- YEARS_BEGINEXPLUATATION_MEDI: double (nullable = true)
|-- YEARS_BUILD_MEDI: double (nullable = true)
|-- COMMONAREA_MEDI: double (nullable = true)
|-- ELEVATORS_MEDI: double (nullable = true)
|-- ENTRANCES_MEDI: double (nullable = true)
|-- FLOORSMAX_MEDI: double (nullable = true)
|-- FLOORSMIN_MEDI: double (nullable = true)
|-- LANDAREA_MEDI: double (nullable = true)
```

```
|-- LIVINGAPARTMENTS_MEDI: double (nullable = true)
|-- LIVINGAREA_MEDI: double (nullable = true)
|-- NONLIVINGAPARTMENTS_MEDI: double (nullable = true)
|-- NONLIVINGAREA_MEDI: double (nullable = true)
|-- FONDKAPREMONT_MODE: string (nullable = true)
|-- HOUSETYPE_MODE: string (nullable = true)
|-- TOTALAREA_MODE: double (nullable = true)
|-- WALLSMATERIAL_MODE: string (nullable = true)
|-- EMERGENCYSTATE_MODE: string (nullable = true)
|-- OBS_30_CNT_SOCIAL_CIRCLE: double (nullable = true)
|-- DEF_30_CNT_SOCIAL_CIRCLE: double (nullable = true)
|-- OBS_60_CNT_SOCIAL_CIRCLE: double (nullable = true)
|-- DEF_60_CNT_SOCIAL_CIRCLE: double (nullable = true)
|-- DAYS_LAST_PHONE_CHANGE: double (nullable = true)
|-- FLAG_DOCUMENT_2: integer (nullable = true)
|-- FLAG_DOCUMENT_3: integer (nullable = true)
|-- FLAG_DOCUMENT_4: integer (nullable = true)
|-- FLAG_DOCUMENT_5: integer (nullable = true)
|-- FLAG_DOCUMENT_6: integer (nullable = true)
|-- FLAG_DOCUMENT_7: integer (nullable = true)
|-- FLAG_DOCUMENT_8: integer (nullable = true)
|-- FLAG_DOCUMENT_9: integer (nullable = true)
|-- FLAG_DOCUMENT_10: integer (nullable = true)
|-- FLAG_DOCUMENT_11: integer (nullable = true)
|-- FLAG_DOCUMENT_12: integer (nullable = true)
|-- FLAG_DOCUMENT_13: integer (nullable = true)
|-- FLAG_DOCUMENT_14: integer (nullable = true)
|-- FLAG_DOCUMENT_15: integer (nullable = true)
|-- FLAG_DOCUMENT_16: integer (nullable = true)
|-- FLAG_DOCUMENT_17: integer (nullable = true)
|-- FLAG_DOCUMENT_18: integer (nullable = true)
|-- FLAG_DOCUMENT_19: integer (nullable = true)
|-- FLAG_DOCUMENT_20: integer (nullable = true)
|-- FLAG_DOCUMENT_21: integer (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_HOUR: double (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_DAY: double (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_WEEK: double (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_MON: double (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_QRT: double (nullable = true)
|-- AMT_REQ_CREDIT_BUREAU_YEAR: double (nullable = true)
```

- printSchema() only shows us the column names and its data type. we are going to drop the SK_ID_CURR column, rename the “TARGET” column to “label” and see the distribution of our target variable:

2. Analysis of the Data

```
In [ ]: # Sk_ID_Curr is the id column which we dont need it in the process #so we get rid of it. and we rename the name of our
# target variable to "label"
drop_col = ['SK_ID_CURR']
new_df = new_df.select([column for column in new_df.columns if column not in drop_col])
new_df = new_df.withColumnRenamed('TARGET', 'label')
new_df.groupby('label').count().toPandas()
```

Out[]:

	label	count
0	1	24825
1	0	282686

- Dataset is highly imbalanced

```
In [ ]: # Let's have a look at the distribution of our target variable:
# to make it look better, we first convert our spark df to a Pandas
# Import the updated library
import pandas as pd
# Rest of your code
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
df_pd = new_df.toPandas()
print(len(df_pd))
plt.figure(figsize=(6,3))
sns.countplot(x='label', data=df_pd, order=df_pd['label'].value_counts().index)
```

307511


```
d:\Personal Files\Projects\Credit_Risk_Modelling_using_PySpark\venv\lib\site-packages\seaborn\_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
```

```
if pd.api.types.is_categorical_dtype(vector):
```

```
d:\Personal Files\Projects\Credit_Risk_Modelling_using_PySpark\venv\lib\site-packages\seaborn\_oldcore.py:1498: FutureWarning: is_categorical_dtype is deprecated and will be removed in a future version. Use isinstance(dtype, CategoricalDtype) instead
```

```
if pd.api.types.is_categorical_dtype(vector):
```

```
Out[ ]: <Axes: xlabel='label', ylabel='count'>
```



Data Wrangling can be summarized in the following points:

1. Understanding Dataset Structure: Begin by gaining insights into the dataset's general structure.
2. Identifying Feature Types: Determine the count of Categorical and Numerical features within the dataset.
3. Handling Missing Values: Develop a function to extract crucial information regarding missing values in the dataset.

Identifying Categorical & Numerical features


```
In [ ]: # now Let's see how many categorical and numerical features we have:
cat_cols = [item[0] for item in new_df.dtypes if item[1].startswith('string')]
print(str(len(cat_cols)) + ' categorical features')
num_cols = [item[0] for item in new_df.dtypes if item[1].startswith('int') | item[1].startswith('double')][1:]
print(str(len(num_cols)) + ' numerical features')

16 categorical features
104 numerical features
```

Identifying missing info

```

In [ ]: # we use the below function to find more information about the #missing values
def info_missing_table(df_pd):
    """Input pandas dataframe and Return columns with missing value and percentage"""
    mis_val = df_pd.isnull().sum() #count total of null in each columns in dataframe
    #count percentage of null in each columns
    mis_val_percent = 100 * df_pd.isnull().sum() / len(df_pd)
    mis_val_table = pd.concat([mis_val, mis_val_percent], axis=1)
    #join to left (as column) between mis_val and mis_val_percent
    mis_val_table_ren_columns = mis_val_table.rename(
        columns = {0 : 'Missing Values', 1 : '% of Total Values'})
    #rename columns in table
    mis_val_table_ren_columns = mis_val_table_ren_columns[
        mis_val_table_ren_columns.iloc[:,1] != 0].sort_values('% of Total Values', ascending=False).round(1)

    print ("Your selected dataframe has " + str(df_pd.shape[1]) + " columns.\n"      #.shape[1] : just view total column
s in dataframe
    "There are " + str(mis_val_table_ren_columns.shape[0]) +
    " columns that have missing values.") #.shape[0] : just view total rows in dataframe
    return mis_val_table_ren_columns

missings = info_missing_table(df_pd)
missings

```

Your selected dataframe has 121 columns.
There are 67 columns that have missing values.

Out[]:

	Missing Values	% of Total Values
COMMONAREA_MEDI	214865	69.9
COMMONAREA_AVG	214865	69.9
COMMONAREA_MODE	214865	69.9
NONLIVINGAPARTMENTS_MEDI	213514	69.4
NONLIVINGAPARTMENTS_MODE	213514	69.4
...
EXT_SOURCE_2	660	0.2
AMT_GOODS_PRICE	278	0.1
AMT_ANNUITY	12	0.0
CNT_FAM_MEMBERS	2	0.0
DAYS_LAST_PHONE_CHANGE	1	0.0

67 rows × 2 columns

- There are 67 columns out of 121 that has missing values in them. and it doesn't show all of them in the image, but overall, most of these 67 columns have more than 50 percent missing values. so we are dealing with a lot of missing values. we are going to fill the numerical missing values with the average of each column and the categorical missing values with the most frequent category of each column.

```
In [ ]: # so this function deals with the a spark dataframe directly to find more about the missing values
def count_missings(spark_df):
    null_counts = []
    for col in spark_df.dtypes:
        cname = col[0]
        ctype = col[1]
        nulls = spark_df.where( spark_df[cname].isNull()).count() #check count of null in column name
        result = tuple([cname, nulls]) #new tuple, (column name, null count)
        null_counts.append(result) #put the new tuple in our result list
    null_counts=[(x,y) for (x,y) in null_counts if y!=0] #view just columns that have missing values
    return null_counts
```

```
In [ ]: # counts =missings.index
# miss_counts = []
# for column, miss_count in zip(missings.index, missings['Missing Values']):
#     miss_counts.append((column, miss_count))

miss_counts = count_missings(new_df)
miss_counts
```

```
Out[ ]: [('AMT_ANNUITY', 12),
 ('AMT_GOODS_PRICE', 278),
 ('NAME_TYPE_SUITE', 1292),
 ('OWN_CAR_AGE', 202929),
 ('OCCUPATION_TYPE', 96391),
 ('CNT_FAM_MEMBERS', 2),
 ('EXT_SOURCE_1', 173378),
 ('EXT_SOURCE_2', 660),
 ('EXT_SOURCE_3', 60965),
 ('APARTMENTS_AVG', 156061),
 ('BASEMENTAREA_AVG', 179943),
 ('YEARS_BEGINEXPLUATATION_AVG', 150007),
 ('YEARS_BUILD_AVG', 204488),
 ('COMMONAREA_AVG', 214865),
 ('ELEVATORS_AVG', 163891),
 ('ENTRANCES_AVG', 154828),
 ('FLOORSMAX_AVG', 153020),
 ('FLOORSMIN_AVG', 208642),
 ('LANDAREA_AVG', 182590),
 ('LIVINGAPARTMENTS_AVG', 210199),
 ('LIVINGAREA_AVG', 154350),
 ('NONLIVINGAPARTMENTS_AVG', 213514),
 ('NONLIVINGAREA_AVG', 169682),
 ('APARTMENTS_MODE', 156061),
 ('BASEMENTAREA_MODE', 179943),
 ('YEARS_BEGINEXPLUATATION_MODE', 150007),
 ('YEARS_BUILD_MODE', 204488),
 ('COMMONAREA_MODE', 214865),
 ('ELEVATORS_MODE', 163891),
 ('ENTRANCES_MODE', 154828),
 ('FLOORSMAX_MODE', 153020),
 ('FLOORSMIN_MODE', 208642),
 ('LANDAREA_MODE', 182590),
 ('LIVINGAPARTMENTS_MODE', 210199),
 ('LIVINGAREA_MODE', 154350),
 ('NONLIVINGAPARTMENTS_MODE', 213514),
 ('NONLIVINGAREA_MODE', 169682),
 ('APARTMENTS_MEDI', 156061),
 ('BASEMENTAREA_MEDI', 179943),
 ('YEARS_BEGINEXPLUATATION_MEDI', 150007),
 ('YEARS_BUILD_MEDI', 204488),
```

```
('COMMONAREA_MEDI', 214865),
('ELEVATORS_MEDI', 163891),
('ENTRANCES_MEDI', 154828),
('FLOORSMAX_MEDI', 153020),
('FLOORSMIN_MEDI', 208642),
('LANDAREA_MEDI', 182590),
('LIVINGAPARTMENTS_MEDI', 210199),
('LIVINGAREA_MEDI', 154350),
('NONLIVINGAPARTMENTS_MEDI', 213514),
('NONLIVINGAREA_MEDI', 169682),
('FONDKAPREMONT_MODE', 210295),
('HOUSETYPE_MODE', 154297),
('TOTALAREA_MODE', 148431),
('WALLSMATERIAL_MODE', 156341),
('EMERGENCYSTATE_MODE', 145755),
('OBS_30_CNT_SOCIAL_CIRCLE', 1021),
('DEF_30_CNT_SOCIAL_CIRCLE', 1021),
('OBS_60_CNT_SOCIAL_CIRCLE', 1021),
('DEF_60_CNT_SOCIAL_CIRCLE', 1021),
('DAYS_LAST_PHONE_CHANGE', 1),
('AMT_REQ_CREDIT_BUREAU_HOUR', 41519),
('AMT_REQ_CREDIT_BUREAU_DAY', 41519),
('AMT_REQ_CREDIT_BUREAU_WEEK', 41519),
('AMT_REQ_CREDIT_BUREAU_MON', 41519),
('AMT_REQ_CREDIT_BUREAU_QRT', 41519),
('AMT_REQ_CREDIT_BUREAU_YEAR', 41519)]
```

3. Data Pre-processing

Separate categorical and numerical columns with missing values

```
In [ ]: # here we separate missing columns in our new_df based on #categorical and numerical types
list_cols_miss=[x[0] for x in miss_counts]
list_cols_miss
df_miss= new_df.select(*list_cols_miss)
# categorical columns
catcolumns_miss=[item[0] for item in df_miss.dtypes if item[1].startswith('string')] #will select name of column with
string data type
print("categorical columns_miss:", catcolumns_miss)
### numerical columns
numcolumns_miss = [item[0] for item in df_miss.dtypes if item[1].startswith('int') | item[1].startswith('double')] #wi
ll select name of column with integer or double data type
print("numerical columns_miss:", numcolumns_miss)
```

```
categorical columns_miss: ['NAME_TYPE_SUITE', 'OCCUPATION_TYPE', 'FONDKAPREMONT_MODE', 'HOUSETYPE_MODE', 'WALLSMATERI
AL_MODE', 'EMERGENCYSTATE_MODE']
numerical columns_miss: ['AMT_ANNUITY', 'AMT_GOODS_PRICE', 'OWN_CAR_AGE', 'CNT_FAM_MEMBERS', 'EXT_SOURCE_1', 'EXT_SOU
RCE_2', 'EXT_SOURCE_3', 'APARTMENTS_AVG', 'BASEMENTAREA_AVG', 'YEARS_BEGINEXPLUATATION_AVG', 'YEARS_BUILD_AVG', 'COMM
ONAREA_AVG', 'ELEVATORS_AVG', 'ENTRANCES_AVG', 'FLOORSMAX_AVG', 'FLOORSMIN_AVG', 'LANDAREA_AVG', 'LIVINGAPARTMENTS_AV
G', 'LIVINGAREA_AVG', 'NONLIVINGAPARTMENTS_AVG', 'NONLIVINGAREA_AVG', 'APARTMENTS_MODE', 'BASEMENTAREA_MODE', 'YEARS_
BEGINEXPLUATATION_MODE', 'YEARS_BUILD_MODE', 'COMMONAREA_MODE', 'ELEVATORS_MODE', 'ENTRANCES_MODE', 'FLOORSMAX_MODE',
'FLOORSMIN_MODE', 'LANDAREA_MODE', 'LIVINGAPARTMENTS_MODE', 'LIVINGAREA_MODE', 'NONLIVINGAPARTMENTS_MODE', 'NONLIVING
AREA_MODE', 'APARTMENTS_MEDI', 'BASEMENTAREA_MEDI', 'YEARS_BEGINEXPLUATATION_MEDI', 'YEARS_BUILD_MEDI', 'COMMONAREA_M
EDI', 'ELEVATORS_MEDI', 'ENTRANCES_MEDI', 'FLOORSMAX_MEDI', 'FLOORSMIN_MEDI', 'LANDAREA_MEDI', 'LIVINGAPARTMENTS_MED
I', 'LIVINGAREA_MEDI', 'NONLIVINGAPARTMENTS_MEDI', 'NONLIVINGAREA_MEDI', 'TOTALAREA_MODE', 'OBS_30_CNT_SOCIAL_CIRCL
E', 'DEF_30_CNT_SOCIAL_CIRCLE', 'OBS_60_CNT_SOCIAL_CIRCLE', 'DEF_60_CNT_SOCIAL_CIRCLE', 'DAYS_LAST_PHONE_CHANGE', 'AM
T_REQ_CREDIT_BUREAU_HOUR', 'AMT_REQ_CREDIT_BUREAU_DAY', 'AMT_REQ_CREDIT_BUREAU_WEEK', 'AMT_REQ_CREDIT_BUREAU_MON', 'A
MT_REQ_CREDIT_BUREAU_QRT', 'AMT_REQ_CREDIT_BUREAU_YEAR']
```

Fill the missing values


```
In [ ]: # now that we have seperated the columns based on categorical and #numerical types, we will fill the missing categoric
L
# values with the most frequent category
from pyspark.sql.functions import rank,sum,col
df_Nomiss=new_df.na.drop()
for x in catcolumns_miss:
    mode=df_Nomiss.groupBy(x).count().sort(col("count").desc()).collect()[0][0]
    print(x, mode) #print name of columns and it's most categories
    new_df = new_df.na.fill({x:mode})
# and we fill the missing numerical values with the average of each #column
from pyspark.sql.functions import mean, round
for i in numcolumns_miss:
    meanvalue = new_df.select(round(mean(i))).collect()[0][0]
    print(i, meanvalue)
    new_df=new_df.na.fill({i:meanvalue})
```

NAME_TYPE_SUITE Unaccompanied
OCCUPATION_TYPE Laborers
FONDKAPREMONT_MODE reg oper account
HOUSETYPE_MODE block of flats
WALLSMATERIAL_MODE Panel
EMERGENCYSTATE_MODE No
AMT_ANNUITY 27109.0
AMT_GOODS_PRICE 538396.0
OWN_CAR_AGE 12.0
CNT_FAM_MEMBERS 2.0
EXT_SOURCE_1 1.0
EXT_SOURCE_2 1.0
EXT_SOURCE_3 1.0
APARTMENTS_AVG 0.0
BASEMENTAREA_AVG 0.0
YEARS_BEGINEXPLUATATION_AVG 1.0
YEARS_BUILD_AVG 1.0
COMMONAREA_AVG 0.0
ELEVATORS_AVG 0.0
ENTRANCES_AVG 0.0
FLOORSMAX_AVG 0.0
FLOORSMIN_AVG 0.0
LANDAREA_AVG 0.0
LIVINGAPARTMENTS_AVG 0.0
LIVINGAREA_AVG 0.0
NONLIVINGAPARTMENTS_AVG 0.0
NONLIVINGAREA_AVG 0.0
APARTMENTS_MODE 0.0
BASEMENTAREA_MODE 0.0
YEARS_BEGINEXPLUATATION_MODE 1.0
YEARS_BUILD_MODE 1.0
COMMONAREA_MODE 0.0
ELEVATORS_MODE 0.0
ENTRANCES_MODE 0.0
FLOORSMAX_MODE 0.0
FLOORSMIN_MODE 0.0
LANDAREA_MODE 0.0
LIVINGAPARTMENTS_MODE 0.0
LIVINGAREA_MODE 0.0
NONLIVINGAPARTMENTS_MODE 0.0
NONLIVINGAREA_MODE 0.0

APARTMENTS_MEDI 0.0
BASEMENTAREA_MEDI 0.0
YEARS_BEGINEXPLUATATION_MEDI 1.0
YEARS_BUILD_MEDI 1.0
COMMONAREA_MEDI 0.0
ELEVATORS_MEDI 0.0
ENTRANCES_MEDI 0.0
FLOORSMAX_MEDI 0.0
FLOORSMIN_MEDI 0.0
LANDAREA_MEDI 0.0
LIVINGAPARTMENTS_MEDI 0.0
LIVINGAREA_MEDI 0.0
NONLIVINGAPARTMENTS_MEDI 0.0
NONLIVINGAREA_MEDI 0.0
TOTALAREA_MODE 0.0
OBS_30_CNT_SOCIAL_CIRCLE 1.0
DEF_30_CNT_SOCIAL_CIRCLE 0.0
OBS_60_CNT_SOCIAL_CIRCLE 1.0
DEF_60_CNT_SOCIAL_CIRCLE 0.0
DAYS_LAST_PHONE_CHANGE -963.0
AMT_REQ_CREDIT_BUREAU_HOUR 0.0
AMT_REQ_CREDIT_BUREAU_DAY 0.0
AMT_REQ_CREDIT_BUREAU_WEEK 0.0
AMT_REQ_CREDIT_BUREAU_MON 0.0
AMT_REQ_CREDIT_BUREAU_QRT 0.0
AMT_REQ_CREDIT_BUREAU_YEAR 2.0

- The dataset no longer contains missing values.

Handling Imbalanced Classess

- The next step is addressing the issue of imbalanced classes.
- Various methods can be used to mitigate this problem.
- One approach involves under-sampling the majority class or over-sampling the minority class to achieve a more balanced outcome.
- Another approach is assigning weights to each class to penalize the majority class with lower weights and boost the minority class with higher weights.
- To implement this, a new column called "weights" is created in the dataset.
- The weights are assigned based on the inverse ratio of each class.

```
In [ ]: # adding the new column weights and fill it with ratios
from pyspark.sql.functions import when
ratio = 0.91
def weight_balance(labels):
    return when(labels == 1, ratio).otherwise(1*(1-ratio))
new_df = new_df.withColumn('weights', weight_balance(col('label')))
```

```
In [ ]: new_df.show(5)
```

	label	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUI
	TY	AMT_GOODS_PRICE	NAME_TYPE_SUITE	NAME_INCOME_TYPE	NAME_EDUCATION_TYPE	NAME_FAMILY_STATUS	NAME_HOUSING_TYPE	REGIO	N_POPULATION_RELATIVE
	DAYS_BIRTH	DAYS_EMPLOYED	DAYS_REGISTRATION	DAYS_ID_PUBLISH	OWN_CAR_AGE	FLAG_MOBIL	FLAG_EMP_PHON	E	FLAG_WORK_PHONE
	FLAG_CONT_MOBILE	FLAG_PHONE	FLAG_EMAIL	OCCUPATION_TYPE	CNT_FAM_MEMBERS	REGION_RATING_CLIENT	REGION_RATING_CLIENT_W_CITY	WEEKDAY_APPR_PROCESS_START	HOUR_APPR_PROCESS_START
	REG_REGION_NOT_LIVE_REGION	REG_REGION_NOT_WOR	K_REGION	LIVE_REGION_NOT_WORK_REGION	REG_CITY_NOT_LIVE_CITY	REG_CITY_NOT_WORK_CITY	LIVE_CITY_NOT_WORK_CITY	ORGANIZ	ATION_TYPE
	EXT_SOURCE_1	EXT_SOURCE_2	EXT_SOURCE_3	APARTMENTS_AVG	BASEMENTAREA_AVG	YEARS_BEGINEXPL	UATATION_AVG	YEARS_BUILD_AVG	COMMONAREA_AVG
	ELEVATORS_AVG	ENTRANCES_AVG	FLOORSMAX_AVG	FLOORSMIN_AVG	LANDAREA_AVG	L	VIVINGAPARTMENTS_AVG	LIVINGAREA_AVG	NONLIVINGAPARTMENTS_AVG
	NONLIVINGAREA_AVG	APARTMENTS_MODE	BASEMENTAREA_MODE	YEARS_BEGINEXPLUATATION_MODE	YEARS_BUILD_MODE	COMMONAREA_MODE	ELEVATORS_MODE	ENTRANCES_MODE	FLOORSMAX_MODE
	FLOORSMIN_MODE	L	ANDAREA_MODE	LIVINGAPARTMENTS_MODE	LIVINGAREA_MODE	NONLIVINGAPARTMENTS_MODE	NONLIVINGAREA_MODE	APARTMENTS_MEDI	BASEME
	NTAREA_MEDI	YEARS_BEGINEXPLUATATION_MEDI	YEARS_BUILD_MEDI	COMMONAREA_MEDI	ELEVATORS_MEDI	ENTRANCES_MEDI	FLOORSMAX_MED	I	FLOORSMIN_MEDI
	LANDAREA_MEDI	LIVINGAPARTMENTS_MEDI	LIVINGAREA_MEDI	NONLIVINGAPARTMENTS_MEDI	NONLIVINGAREA_MEDI	FOND	KAPREMONT_MODE	HOUSETYPE_MODE	TOTALAREA_MODE
	WALLSMATERIAL_MODE	EMERGENCYSTATE_MODE	OBS_30_CNT_SOCIAL_CIRCLE	DEF_30_C	NT_SOCIAL_CIRCLE	OBS_60_CNT_SOCIAL_CIRCLE	DEF_60_CNT_SOCIAL_CIRCLE	DAYS_LAST_PHONE_CHANGE	FLAG_DOCUMENT_2
	FLAG_DOCUMENT_3	FLAG_DOCUMENT_4	FLAG_DOCUMENT_5	FLAG_DOCUMENT_6	FLAG_DOCUMENT_7	FLAG_DOCUMENT_8	FLAG_DOCUMENT_9	FLAG_DOCUMENT_10	FLAG_DOCUMENT_11
	FLAG_DOCUMENT_12	FLAG_DOCUMENT_13	FLAG_DOCUMENT_14	FLAG_DOCUMENT_15	FLAG_DOCUMENT_16	FLAG_DOCUMENT_17	FLAG_DOCUMENT_18	FLAG_DOCUMENT_19	FLAG_DOCUMENT_20
	FLAG_DOCUMENT_21	AMT_REQ_CREDIT_BUREAU_HOUR	AMT_REQ_CREDIT_BURE	AU_DAY	AMT_REQ_CREDIT_BUREAU_WEEK	AMT_REQ_CREDIT_BUREAU_MON	AMT_REQ_CREDIT_BUREAU_QRT	AMT_REQ_CREDIT_BUREAU_YEAR	weights

1	Cash loans	M	N	Y	0	202500.0	406597.5	2470
0.5	351000.0	Unaccompanied	Working	Secondary / secon...	Single / not married	House / apartment		
0.018801	-9461	-637	-3648.0	-2120	12.0	1	1	
0	1	1	0	Laborers	1.0	2		
2	WEDNESDAY		10		0		0	
0	0		0		0	Business Entity T...	0.08303696739132256	0.26
29485927471776	0.13937578009978951		0.0247	0.0369		0.9722		0.6192
0.0143	0.0	0.069	0.0833	0.125	0.0369	0.0202		0.019
0.0	0.0	0.0252	0.0383		0.9722	0.6341		0.0144
	0.0	0.069	0.0833	0.125	0.0377	0.022		0.0198
0.0	0.0	0.025	0.0369		0.9722	0.6243		0.014
4	0.0	0.069	0.0833	0.125	0.0375	0.0205		0.0193
0.0	0.0	reg oper account	block of flats		0.0149	Stone, brick		No
2.0		2.0	2.0		2.0	-1134.0		0
	1	0	0	0	0	0		0
0	0	0	0	0	0	0		0
0	0	0	0	0	0	0.0		0
0.0		0.0		0.0		0.0		1.0
0.91								
	Cash loans	F	N	N	0	270000.0	1293502.5	3569
8.5	1129500.0	Family	State servant	Higher education		Married	House / apartment	
0.0035409999999999999	-16765	-1188	-1186.0	-291	12.0	1		1
	0	1	1	0	Core staff	2.0		1

1		MONDAY		11		0		0
0		0		0		0	School	0.3112673113812225 0.62
22457752555098		1.0		0.0959		0.0529		0.9851 0.7959999999999999
0.0605		0.08		0.0345		0.2917		0.3333
0.0039		0.0098		0.0924		0.0538		0.9851
497		0.0806		0.0345		0.2917		0.3333
0.0		0.0		0.0968		0.0529		0.9851
8		0.08		0.0345		0.2917		0.3333
0.0039		0.01		reg oper account		block of flats		0.0714
1.0		0.0		1.0		0.0		-828.0
		1		0		0		0
0		0		0		0		0
0		0		0		0		0.0
0.0		0.0		0.0		0.0		0.0 0.08999
999999999997								
		0		Revolving loans		M		Y
0.0		135000.0		Unaccompanied		Working		Secondary / secon...
0.010032		-19046		-225		-4260.0		-2531
1		1		1		0		Laborers
2				MONDAY		9		0
0		0		0		0		Government
59120833904428		0.7295666907060153		0.0		0.0		1.0
0.0		0.0		0.0		0.0		0.0
0.0		0.0		0.0		0.0		1.0
		0.0		0.0		0.0		0.0
0.0		0.0		0.0		0.0		1.0
0		0.0		0.0		0.0		0.0
0.0		0.0		0.0		0.0		1.0
0.0		0.0		0.0		0.0		0.0
0.0		0.0		reg oper account		block of flats		0.0
0.0		0.0		0.0		0.0		Panel
		0		0		0		0
0		0		0		0		0
0		0		0		0		0.0
0.0		0.0		0.0		0.0		0.0 0.08999
999999999997								
		0		Cash loans		F		N
6.5		297000.0		Unaccompanied		Working		Secondary / secon...
0.008019		-19005		-3039		-9833.0		-2437
0		1		0		0		Laborers
2				WEDNESDAY		17		0
0		0		0		0		Business Entity T...
04416904014653		1.0		0.0		0.0		1.0
0.0		0.0		0.0		0.0		0.0


```

-+-+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
---+-----+-----+-----+-----+-----+-----+-----+-----+-----+
+-+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-+-+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 5 rows

```

Feature Engineering

- The next step in the process is Feature Engineering.
- pySpark simplifies feature extraction, making it easier.
- The steps involved in feature extraction are as follows:
 1. Apply StringIndexer() to assign indices to each category in our categorical columns.
 2. Use OneHotEncoderEstimator() to convert categorical columns into one-hot encoded vectors.
 3. Apply VectorAssembler() to create a feature vector that combines all categorical and numerical features, referred to as "features."

```

In [ ]: # We use the OneHotEncoderEstimator from MLlib in Spark to convert each categorical feature into one-hot vectors.
# Next, we use VectorAssembler to combine the resulting one-hot vectors and the rest of the numerical features into a
single vector column.
# We append every step of the process in a stages array.

from pyspark.ml.feature import OneHotEncoder, StringIndexer, VectorAssembler

stages = []

for categoricalCol in cat_cols:
    stringIndexer = StringIndexer(inputCol=categoricalCol, outputCol=categoricalCol + 'Index')
    encoder = OneHotEncoder(inputCols=[stringIndexer.getOutputCol()], outputCols=[categoricalCol + "classVec"])
    stages += [stringIndexer, encoder]

assemblerInputs = [c + "classVec" for c in cat_cols] + num_cols
assembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")
stages += [assembler]

```

1. We will now create a pipeline.
2. This pipeline will perform a sequence of transformations.
3. The purpose is to apply all these transformations at once.
4. This approach simplifies the process of applying multiple transformations in a sequence.

```
In [ ]: # we use a pipeline to apply all the stages of transformation
from pyspark.ml import Pipeline
cols = new_df.columns
pipeline = Pipeline(stages = stages)
pipelineModel = pipeline.fit(new_df)
new_df = pipelineModel.transform(new_df)
selectedCols = ['features']+cols
new_df = new_df.select(selectedCols)
pd.DataFrame(new_df.take(5), columns=new_df.columns)
```

Out[]:

	features	label	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME
0	(1.0, 0.0, 1.0, 1.0, 1.0, 1.0, 0.0, 0.0, 0.0, ...	1	Cash loans	M	N	Y	0	202500.0
1	(1.0, 1.0, 0.0, 1.0, 0.0, 0.0, 1.0, 0.0, 0.0, ...	0	Cash loans	F	N	N	0	270000.0
2	(0.0, 0.0, 1.0, 0.0, 1.0, 1.0, 0.0, 0.0, 0.0, ...	0	Revolving loans	M	Y	Y	0	67500.0
3	(1.0, 1.0, 0.0, 1.0, 1.0, 1.0, 0.0, 0.0, 0.0, ...	0	Cash loans	F	N	Y	0	135000.0
4	(1.0, 0.0, 1.0, 1.0, 1.0, 1.0, 0.0, 0.0, 0.0, ...	0	Cash loans	M	N	Y	0	121500.0

5 rows × 123 columns

4. Model Building

Training and Hyper-parameter Tuning

- Split the dataset into training and testing sets for training.
- Begin training with Logistic Regression.
- Choose Logistic Regression for its performance in binary classification problems.

Train-Test Split

```
In [ ]: # split the data into trainign and testin sets
train, test = new_df.randomSplit([0.80, 0.20], seed = 42)
print(train.count())
print(test.count())
```

246240

61271

Logistics Regression

```
In [ ]: # first we check how LogisticRegression perform
from pyspark.ml.classification import LogisticRegression
LR = LogisticRegression(featuresCol = 'features', labelCol = 'label', maxIter=15)
LR_model = LR.fit(train)
```

Gradient Boosting Trees (GBT)

```
In [ ]: # next we checkout gradient boosting trees
        from pyspark.ml.classification import GBTClassifier
        gbt = GBTClassifier(maxIter=15)
        GBT_Model = gbt.fit(train)
```

- Achieved a significantly improved result of 0.732 using GBT.

Hyperparameter Tuning - Gradient Boosting Trees (GBT)

- Final strategy includes hyper-parameter tuning through grid search.
- Planning to follow up with cross-validation to further enhance GBT performance.

```
In [ ]: from pyspark.ml.tuning import ParamGridBuilder, CrossValidator
        from pyspark.ml.evaluation import BinaryClassificationEvaluator # Import BinaryClassificationEvaluator

        paramGrid = (ParamGridBuilder().addGrid(gbt.maxDepth, [2, 4, 6]).addGrid(gbt.maxBins, [20, 30]).addGrid(gbt.maxIter,
        [10, 15]).build())
        cv = CrossValidator(estimator=gbt, estimatorParamMaps=paramGrid, evaluator = BinaryClassificationEvaluator(), numFolds
        =5)
        # Run cross validations.
        cvModel = cv.fit(train)
```

- The result showed a slight improvement.
- There is room for further enhancement by experimenting with hyper-parameter tuning.

5. Model Evaluation

1. We will create an ROC curve for the training data.
2. The purpose is to assess the performance of Logistic Regression.
3. We will utilize the Area Under the ROC Curve (AUC-ROC).
4. AUC-ROC is a common metric for evaluating binary classification models.

```
In [ ]: from pyspark.ml.evaluation import BinaryClassificationEvaluator
evaluator = BinaryClassificationEvaluator()

predictions_LR = LR_model.transform(test)
print("Test_SET Logistic Regression (Area Under ROC): " + str(evaluator.evaluate(predictions_LR, {evaluator.metricName: "areaUnderROC"})))


gbt_predictions = GBT_Model.transform(test)
print("Test_SET GBT (Area Under ROC): " + str(evaluator.evaluate(gbt_predictions, {evaluator.metricName: "areaUnderROC"})))

gbt_cv_predictions = cvModel.transform(test)
print("Test_SET GBT_Tuned (Area Under ROC): " + str(evaluator.evaluate(gbt_cv_predictions, {evaluator.metricName: "areaUnderROC"})))

Test_SET Logistic Regression (Area Under ROC): 0.7215699765501389
Test_SET GBT (Area Under ROC): 0.7298618196435773
Test_SET GBT_Tuned (Area Under ROC): 0.7321482798175072
```



```
In [ ]: #plotting the ROC Curve
trainingSummary = LR_model.summary
roc_LR = trainingSummary.roc.toPandas()
plt.plot(roc_LR['FPR'],roc_LR['TPR'])
plt.ylabel('False Positive Rate')
plt.xlabel('True Positive Rate')
plt.title('ROC Curve - Logistics Regression')
plt.show()
print('Training set ROC: ' + str(trainingSummary.areaUnderROC))
```



Training set ROC: 0.7229899807041205

THE END